



seL4 Developers Day

Beginner sel4 Development

Ihor Kuz, Adrian Danis, Gernot Heiser, Stephen Sherratt

October 2015

www.csiro.au



Agenda and Exercises



Agenda

- seL4 overview
- Building and running
- Finding code and debugging
- System startup
 - Root task and bootinfo
- Untyped and Allocators
- Threads and IPC
- CSpaces, VSpaces
- Processes
 - Vspace library
 - Sel4utils library

Exercises

- Hello-1:
 - simple hello world
- Hello-2:
 - start a new thread
- Hello-3:
 - IPC between threads
- Hello-4:
 - start a new process
 - IPC between processes
- Platform: ia32 (x86)
 - Bonus: make it work for ARM

seL4 Overview

Building and running a basic seL4 program - Hello World

VirtualBox VM

A screenshot of a VirtualBox VM window titled 'sel4-tutorial-lubuntu15.04_1 [Running]'. The window contains two overlapping windows. The background window is an Emacs editor showing a C source file named 'main.c'. The code includes a license notice, a tag '@TAG(NICTA_BSD)', a comment about the tutorial part, and a main function skeleton with a TODO comment. The foreground window is a terminal window titled 'sel4@sel4-VirtualBox: ~/Projects/sel4-tutorials-sel4'. It shows the output of a build process, including compilation steps for various object files, linking, and a final error message: 'make[1]: *** [hello-1.elf] Error 1' and 'make: *** [hello-1] Error 2'. The terminal prompt is 'sel4@sel4-VirtualBox:~/Projects/sel4-tutorials-sel4\$'.

• VirtualBox VM contents:

- Linux: Lubuntu 15.04
 - login: sel4:sel4
- Prerequisites installed
- Tutorial code
 - In ~/Projects/
- Tutorial slides & seL4 docs

Install and run

- Get Virtualbox (newest: 5.0.4)
- Install it
 - Also get and install the Extension Pack
- Get seL4 tutorial “appliance”
- Import it, Run it

Prerequisites



- Knowledge prereqs (what we expect)
 - Not required:
 - Previous seL4 experience
 - Required:
 - Be a quick learner!
 - C programming
 - Unix environment (editing files, running build tools, etc.)
- Technical prereqs (what you should have working already)
 - Linux (will work on BSD, MacOSX with some extra work, cygwin maybe?)
 - C compilers and build tools (gcc)
 - Python and packages
 - Haskell, cabal, and packages
 - Git and repo
 - qemu

Getting the code



- Using repo (and git)

```
mkdir sel4-tutorials
```

```
cd sel4-tutorials
```

```
repo init -u http://github.com/sel4-projects/sel4-tutorials-manifest  
          -m sel4-tutorials.xml
```

```
repo sync
```

- ***Root directory***

- The directory where you did `repo init`

- Exercises

- In `apps/hello-*`
- Contain program skeletons
 - Interesting bits replaced with “TODO” comments
 - Includes lots of hints
- Goal is to add code for the TODOs to create complete programs

Working Directory



- What's there?
 - Config and build files: Kconfig, Kbuild, Makefile
 - kernel/ - from <https://github.com/seL4/seL4>
 - libs/ - from https://github.com/seL4/libs*
 - projects/
 - sel4-tutorials/ - from <https://github.com/seL4-projects/sel4-tutorials>
 - docs/, run-arm.sh, run-ia32.sh
 - apps/
 - hello-*/
 - src/*.c, Kbuild, Kconfig, Makefile
 - configs/
 - *_defconfig
 - tools/
 - - build system (.mk files) and build tools

Documentation



- seL4 manual
 - <https://sel4.systems/Info/Docs/seL4-manual.pdf>
- CAmkES manual
 - <https://github.com/seL4/camkes-tool/blob/master/docs/index.md>
- Code
 - seL4 API
 - <https://github.com/seL4/seL4/tree/master/libsel4>
 - kernel/libsel4
 - Libraries
 - https://github.com/seL4/lib*
 - libs/
- These slides

Configuration



- Based on the Linux Kernel Build System: Kbuild
 - `Kconfig`: define config symbols and their attributes
 - `Kbuild`: defines dependencies between modules
 - `Makefile`: make rules (uses many `.mk` files from `tools/common`)
 - `.config`: stores each config symbol's selected value
- `make menuconfig`
 - Curses-based interface, set various config options
 - Writes a `.config` file
- `configs/`
 - Default configurations (copies of `.config` files for different configurations)
 - E.g.: `ia32_hello-1_defconfig`
- `make ...defconfig`
 - E.g. `make ia32_hello-1_defconfig`
 - Copies `configs/ia32_hello-1_defconfig` to `.config`

Building and Running



- `make`
 - Builds required modules based on `.config`
 - Compilation occurs outside of source directory
 - `build/`: generated files and results of compilation
 - `stage/`: intermediate results, copies of header files, etc.
 - `images/`
 - Loadable system image files. Example:
 - Kernel: `kernel-ia32-pc99`
 - User: `hello-0-image-ia32-pc99`
- Run using `qemu`: system emulator
 - `qemu-system-i386 -nographic -m 512 -kernel <kernel> -initrd <user>`
- `make clean`, `make mrproper`
- Convenience scripts
 - `run-ia32.sh`, `run-arm.sh`

Hands-on: hello-1



- Goal: Get and build a simple seL4 system
- Get code (if you haven't already)
 - Hint: `mkdir ...; cd ...; repo init ...; repo sync`
- `make ia32_hello-1_defconfig`
 - `make menuconfig`
- `make`
 - Compilation failure!
 - Find and fix TODO 1
- Run:
 - `qemu-system-i386 -nographic -m 512 -kernel images/kernel-ia32-pc99 -initrd images/hello-1-image-ia32-pc99`
 - Quit qemu with: C-A x
- Where's the solution?
 - `projects/sel4-tutorials/solutions/hello-1`

Finding code and Debugging

Looking up code



- Github
 - <http://github.com/sel4>
 - Lookup appropriate repo (sel4, lib*, etc.)
 - Search (github search isn't great)
- Cscope
 - Search C code
 - Run directly: `cscope -R`
 - <Tab> between fields
 - Choose file to view
 - Ctrl-D to exit
- Run cscope in browser
 - Vim:
 - http://cscope.sourceforge.net/cscope_vim_tutorial.html
 - `:cscope f g <function>`, C-t
 - C-\ g, C-<space> g
 - Emacs: `xcscope.el`
 - <https://github.com/dkogan/xcscope.el>
 - C-s s s

Hands-on: looking up code



- Find these functions:
 - `simple_default_init_bootinfo()`
 - `allocman_make_vka()`
 - `vka_alloc_tcb()`
 - `seL4_TCB_Configure()`
- Use:
 - Github
 - Cscope
- What about libsel4?
 - It is generated at build time!
 - `libsel4/include/interfaces/sel4.xml`
 - `build/.../libsel4/include/interfaces/sel4_client.h`

Build detail



- Build detail
 - Increase detail with V=
 - Example:
 - make V=0
 - make V=1
 - make V=2
 - make V=3
 - Useful to find out why a build step fails
 - What is being run
 - What are the arguments (e.g., include paths, or library paths)

Debugging



- VM fault

Caught cap fault in send phase at address 0x0 while trying to handle:

vm fault on data at address 0x0 with status 0x6 in thread 0xffaf9900 "rootserver" **at address 0x80480db**

- use objdump

- `objdump -dS build/x86/pc99/hello-1/hello-1.bin | less`

- look for instruction at address: 80480db

```
printf("hello world\n");
```

```
80480d1:      68 40 c8 04 08      push    $0x804c840
```

```
80480d6:      e8 c8 02 00 00      call    80483a3 <puts>
```

```
    *(char*)0x0 = 'a';
```

```
80480db:      c6 05 00 00 00 00 00 00  movb    $0x0,0x0
```

```
80480e2:      0f 0b                ud2
```

Hands-on: hello-1 with errors



- Goal: Introduce an error into hello-1
- Add a line in main():
 - `* (char*) 0x0 = 'X' ;`
- Build and run it -> VM fault
- Find it
 - Use objdump to find it
- Fix it 😊

seL4 System Startup

- Root task and Bootinfo**
- libsimple**

seL4 system startup



- Image
 - Kernel image
 - User-space image
 - Root task & Cpio file containing elf files
- Boot loader
 - Loads kernel into memory
 - Loads user-space image into memory
 - Starts kernel running
- Kernel startup
 - Kernel creates object (untyped, frames for device memory)
 - Kernel creates root task objects
 - Loads and runs root task
- Root Task
 - Responsible for setting up the rest of the system

BootInfo: Start-up Information



- Kernel creates:
 - root task CSpace, root task VSpace, Root task TCB
 - frames for device memory
 - untyped caps for RAM memory
- All startup objects are available to root task
 - kernel places caps to these objects in root task CSpace
 - kernel needs to tell root task
 - what caps it has
 - what the objects are
- Bootinfo
 - info about all the initial objects and the caps to them

Initial caps



- Some Initial Caps

- `seL4_CapNull = 0, /* null cap */`
- `seL4_CapInitThreadTCB = 1, /* initial thread's TCB cap */`
- `seL4_CapInitThreadCNode = 2, /* root CNode cap */`
- `seL4_CapInitThreadVSpace = 3, /* VSpace cap */`
- `seL4_CapBootInfoFrame = 9, /* bootinfo frame cap */`
- `seL4_CapInitThreadIPCBuffer = 10, /* initial thread's IPC buffer frame cap */`

- Some bootinfo fields

- `seL4_SlotRegion untyped; /* untyped-object caps (untyped caps) */`
- `seL4_Word untypedPaddrList[...]; /* physical address of each untyped cap */`
- `seL4_Uint8 untypedSizeBitsList[...]; /* size (2^n) bytes of each untyped cap */`
- `seL4_Word numDeviceRegions; /* number of device regions */`
- `seL4_DeviceRegion deviceRegions[...]; /* device regions */`

Intro to seL4 Libraries



- Goal:
 - Make seL4 programming less “user-unfriendly”
 - Do a bunch of the hard things for you
- Interfaces vs Implementations
 - *Interface*
 - key datastructs
 - function definitions
 - generic code to facilitate use of interface
 - *Implementation*
 - adds implementation-specific parts to datastructs
 - implements interface functions

Key interfaces and libraries



- Key Interfaces
 - **simple**: access to initial caps
 - **vka**: virtual kernel allocator
 - **vspace**: VSpace management
- Key Libraries
 - **libseL4**: seL4 kernel API
 - **allocman** (vka): allocator manager
 - **sel4utils** (vspace, io operations): higher level concepts
- Other Libraries
 - **muslc**: C library
 - **platsupport**, **sel4platsupport**: device access
 - **utils**, **debug**, **benchmark**: other useful functionality

Dependencies



- simple:
 - libsel4
- vka:
 - libsel4, utils
- allocman:
 - vka, libsel4, sel4utils, vspace, utils
- vspace:
 - vka, libsel4, utils
- sel4utils:
 - simple, vka, vspace, platsupport, sel4utils, utils, elf, cpio
- platsupport:
 - utils
- sel4platsupport:
 - simple, vka, vspace, platsupport, sel4utils, utils

Library: Simple



- Easy way to access initial caps
 - Includes: untyped, device memory, initial CSpace, initial VSpace
- Abstracts over spec of initial caps
 - root task: uses bootinfo
 - user-level task: can use bootinfo or some other format
- Key concepts
 - location of resources, caps to resources
 - acquiring resource without cap
- Interfaces defined
 - Simple
- Implemented by
 - simple-default, simple-stable

Simple: API



- Files

- libsel4simple, `#include <simple/simple.h>`
 - `libs/libsel4simple/include/simple/`
- libse4simple-default, `#include <simple-default/simple-default.h>`
 - `libs/libsel4simple-default/include/simple-default/`

- Datastructs

- `simple_t`

- Functions

- `simple_default_init_bootinfo`
- `simple_print`
- `simple_get_*`: `pd, tcb, cnode, node_size`
- `simple_get_nth_untyped`
- `simple_get_frame_*`: `cap, info, vaddr`

Hands-on: hello-2 (part 1)



- Goal: Initialise a simple and look at bootinfo
- `make ia32_hello-2_defconfig`
- **Edit** `apps/hello-2/src/main.c`
- **Fix TODOs**
 - TODO 1: get bootinfo
 - TODO 2: init simple
 - TODO 3: print out bootinfo
- `make`
- `run`
 - Print out showing what's in bootinfo
 - See if it makes any sense...

seL4 API and libraries

- Starting a new thread**

Untyped and retyping



- Untyped Memory Object
 - region of (RAM) memory
 - must be retyped to another object to use it
 - results in a nested tree from a root untyped to other objects:
- Retyping
 - kernel uses part of untyped's memory region to store a new kernel object
 - can only create an object if you have a cap to a big enough untyped object
 - retype provides user with cap to the new object
- seL4_Untyped_Retype
 - `seL4_Untyped_Retype(seL4_Untyped service, int type, int size_bits, seL4_CNode root, int node_index, int node_depth, int node_offset, int num_objects)`

Allocators



- Allocating objects requires
 - Pool of untyped memory to retype into the new objects
 - CSpace slots to put the cap to the new object into
 - Memory for bookkeeping structures
 - Bookkeeping:
 - What untyped memory is available, what sizes are they, where are their caps?
 - How much of the untyped memory has been used?
 - What CSpace slots are available?
 - Which objects have been created, and which untyped memory was used and in which Cspace slots are their caps stored?
- Allocator
 - Manages the untyped pools, the CSpace slots
 - Takes care of bookkeeping
 - Maybe even allows objects to be freed!

Library: VKA



- VKA: Virtual Kernel Allocator
- Interface for allocating kernel objects
 - abstracts away
 - creation of objects through retyping untyped
 - managing CSpace and book keeping
- Key concepts
 - **vka**: allocator
 - **Objects**: represent kernel objects
 - **CSpace slots**: slot in local cspace where caps can be found
 - **cspace path**: fully qualified capability address
 - **utSPACE**: pool of untyped memory, used to create objects
- Interfaces defined: `vka`
- Implemented by: allocman

VKA: API



- Files
 - libsel4vka, #include <vka/...> : vka.h, object.h
 - `libs/libsel4vka/include/vka/`
- Datastructs
 - `vka_t`: a VKA interface instance
 - `vka_object_t`: VKA representation of a kernel object
 - Contains: `cptr`, `ut` cookie, object type, size
- Functions
 - `vka_alloc_*`: `pd`, `cnode`, `tcb`, `endpoint`. returns `vka_object_t`
 - `vka_cspace_alloc`: allocate an empty slot in the CSpace
 - `vka_cspace_free`: doesn't delete object!
 - `vka_utspace_alloc`: given empty slot create an object and put the cap to it in the slot. returns `ut` cookie
 - `vka_utspace_free`: given `ut` cookie, free object

Library: Allocman



- Allocator Manager
 - implements vka interface
 - framework combining independent CSpace and utspace allocators
 - solves difficult recursion problems in allocation: Black Magic!
- Key concepts
 - resources: allocator needs underlying resources (e.g. untyped)
 - memory pool: needs an initial pool for internal allocations
- Interfaces implemented
 - vka
 - allocman: to add resources after initialisation

Allocman: API



- Files:
 - libseL4allocman, `#include <allocman/...>`: `allocman.h`, `vka.h`, `bootstrap.h`
 - `libs/libseL4allocman/include/allocman`
- Datastructs
 - `allocman_t`
- Functions
 - Bootstrap: create new allocman
 - `bootstrap_use_current_simple`
 - use current CSpace and simple
 - `bootstrap_new_2level_simple`
 - create and switch to a new CSpace
 - `allocman_make_vka`
 - get VKA interface to allocman

Hands-on: hello-2 (part 2)



- Goal: Create an allocator
- Fix TODOs:
 - TODO 4: create an allocator
 - TODO 5: create a vka
- Build and run
 - No new visible output.

seL4 API: Overview



- Key Concepts
 - Kernel Object
 - in-kernel datastruct, only directly accessible by kernel
 - Capability
 - reference to a kernel object
 - allows holder to invoke functions on the objects
 - i.e. ask kernel to do something with the object
 - holder: thread invoking the cap
- Low-level interface for key activities
 - create kernel objects (retype untyped)
 - create and manage caps in a CSpace
 - create and manage VSpace
 - create and manage threads
 - communicate between threads (IPC)

seL4 API: key files



- Files:
 - Libsel4, `#include <sel4/sel4.h>`
 - Includes: `types.h`, `bootinfo.h`, `arch/syscalls.h`, `interfaces/sel4_client.h`
 - `libs/libsel4/`
 - `include/sel4`
 - `arch_include/x86/sel4/`
 - `build/x86/pc99/libsel4/include/interfaces/`
- Generated from
 - `libs/libsel4/include/interfaces/sel4.xml`
- Note:
 - libsel4 actually lives in the kernel
 - `libs/libsel4` is a symlink to `kernel/libsel4`

seL4 Capability



- Kernel-maintained
 - user-level cannot directly access or manipulate a capability
 - capability is stored in CSpace
 - pass CSpace address of cap in system calls
- Datatypes
 - `seL4_CPtr`
 - index into current thread's CSpace root (CNode)
 - this can be tricky....

seL4 API: TCB (Thread Control Block)



- TCB Object:
 - kernel's representation of a thread.
 - contains:
 - Caps: CSpace, VSpace, IPC Buffer Frame
 - Other: IP (instruction pointer), SP (stack pointer), IPC Buffer, Priority
- IPC Buffer
 - buffer used to pass data during IPC
 - 512 byte object, must be wholly in one frame
 - also used for all syscalls
 - passed to TCB as:
 - index to Frame cap in TCB's CSpace
 - address where it is mapped in TCB's VSpace

TCB: API



- **Configure**
 - `seL4_TCB_Configure`: set CSpace, VSpace, IPC buffer, priority
- **Write Registers**
 - `seL4_TCB_ReadRegisters`: retrieve current registers (ip, sp, etc.)
 - `seL4_TCB_WriteRegisters`: set the current registers (ip, sp, etc.)
 - Set of registers is arch-specific
 - Defined in arch-specific struct
- **Resume**
 - `seL4_TCB_Resume`: start thread running at it's current instruction pointer
- **Suspend**
 - `seL4_TCB_Suspend`: stop thread running

Hands-on: hello-2 (part 3)



- Goal: Create a thread
- TODO 6: get CSpace root Cnode
- TODO 7: get our VSpace root page directory
- TODO 8: create a new TCB
- TODO 9: initialise the new TCB
- TODO 10: give the new thread a name
- TODO 11: set instruction pointer
- TODO 12: set stack pointer
- TODO 13: actually write the TCB registers
- TODO 14: start the new thread running
- TODO 15: print something in new thread
- Wow – that’s a lot of effort for a thread!

More seL4 API and libraries

- IPC

IPC in seL4 – Endpoints



- Endpoint Object
 - Formerly: synchronous endpoint object
 - enables synchronous (blocking) communication
 - communicating threads must hold caps to same endpoint
- Endpoint Caps
 - **master cap** (typically receiver): received when creating the endpoint object
 - **derived caps** (senders): minted from master (or other) caps
 - **badge**: identifies specific sender cap
 - **reply cap**: temporary cap allows receiver to reply to sender for two-way communication
- Message Registers
 - Data to be sent is stored in message registers (MR)
 - Stored in machine registers or in IPC buffer

IPC – Endpoints: API



- **Sending and Receiving**
 - `seL4_Send`: send message registers. Blocks if receiver not Waiting.
 - `seL4_Wait`: wait for a send on endpoint. Blocks if no send pending.
 - `seL4_Call`: send and wait in one syscall. Also sends a reply cap.
 - `seL4_Reply`: send a message using reply cap.
 - `seL4_ReplyWait`: send and wait in one syscall. Using reply cap to send.
- **Message Registers**
 - `seL4_GetMR`: retrieve a given message register from IPC buffer.
 - `seL4_SetMR`: set a given message register in IPC buffer.
 - `seL4_GetCap`: retrieve a cap sent in an IPC.
 - `seL4_SetCap`: prepare a cap to send in an IPC.
- **Tag: `seL4_MessageInfo_t`**
 - Label, message length, number of caps, caps unwrapped

seL4 IPC – Notification



- Notification Object
 - Formerly: asynchronous endpoint object
 - allows one thread to send a notification to another
 - notification: asynchronous (non-blocking) message
 - sends limited data (32-bit word)
 - sets (ORs) bits in receiver notification word, not queued
- Notification Object Caps
 - master cap (receiver), derived caps (senders)
- API
 - `seL4_Notify`: set (OR) notification's data word
 - Also ORs sender badge
 - `seL4_Wait`: retrieve (and clear) notification's data word.
 - Blocks if no new notification since last wait

seL4 API: CSpace



- CNode Object
 - consists of slots in which capabilities are stored
 - can also store CNode caps in slots, creates hierarchical CSpace structure
- API
 - insert cap:
 - indirectly: through `seL4_Untyped_Retype`
 - `seL4_CNode_Copy`: copy from one slot into another
 - `seL4_CNode_Mint`: copy but change some cap attributes
 - `seL4_CNode_Move`: remove from one slot and put in another
 - remove cap
 - `seL4_CNode_Delete`: remove cap from cslot, destroy object if last cap
 - `seL4_CNode_Revoke`: delete all child caps (e.g. through copy or retype)
 - `seL4_CNode_Recycle`: revoke and then reset object attributes

Cspace slots and addresses



- CSpace structure
 - Hierarchy of CNodes: 1-level, 2-level, 3-level, ...
- CNode size (radix)
 - 2^{radix} slots
- CSpace Slot (CSlot) address
 - CPtr: 32-bit Word
 - CPtr is resolved based on CSpace structure (see examples)
 - CPtr resolution relative to a **root cnode**
 - **Depth**: how many bits of CPtr to resolve
- Use of CSlot addresses in syscalls:
 - Just CPtr (implicit root = TCB's CNode cap & implicit 32-bit depth)
 - Explicit root (is a 32-bit depth CPtr with implicit root) & CPtr (called index) & explicit depth

Cspace example and addressing

CNode 0: 2^{16}		Cnode 1: 2^{16}		Cnode 2: 2^{16}		Cnode 3: 2^{16}		Cnode 4	
0		0		0		0		0	
1	A	1	W: cnode 2	1		1		1	E
2	U: cnode 1	2	B	2	C	2	D	2	
3	V: cnode 0	3	X: cnode 0	3		3	Z: cnode 4	3	
4		4		4	Y: cnode 3	4		4	
...

- B: CPtr: 0x 0002 0002
- W: CPtr: 0x 0002 0001
- X: CPtr: 0x 0002 0003
- A: CPtr: 0x 0001 **????**
 - CPtr: 0x 0003 0001
 - Or: Root: X index 0x 0003 depth: 16
- U: CPtr: 0x 0002 **????**
 - CPtr: 0x 0003 0002
 - Or: Root: X index: 0x 0002 depth: 16
- C: CPtr: 0x 0002 0001 **0002**
 - Root: W index: 0x 0002 depth: 16
- Y: CPtr: 0x 0002 0001 **0004**
 - Root: W index: 0x 0004 depth 16
- D: root: W index: 0x 0004 0002 depth: 32
- Z: root: W index: 0x 0004 0003 depth: 32
- E: root: W index: 0x 0004 0003 **0001**


CSpace and guards



- CNode caps can also have guards
 - Inspired by *guarded page tables*
 - Enables sparseness
- Guard
 - Fixed prefix of a CNode's address fragment
 - Example:
 - CNode: radix = 8, guard size = 8 bits value = 0x44
 - CNode's address fragment must be 0x44XX
 - CNode: radix = 12, guard size = 4 bits value = 0x8
 - CNode's address fragment must be 0x4XXX
 - 2-level CSpace: CNode 0: r: 8 g: 8,0x44 CNode 1: r: 12 g: 4,0x4
 - CPtrs must be 0x 44XX 4XXX

seL4 API: VSpace

- VSpace:
 - Objects: PageDir (PD), PageTable (PT), Frame
 - represents mapping: virtual address → physical address
 - i.e. abstraction of CPU page table
- Example (on ARM):



seL4 API: Vspace (contd.)



- VSpace-related objects are platform-specific
- Sizes (ARM)
 - PD: 16KiB, 4 byte slots
 - PT: 1KiB, 4 byte slots
 - Frame: 4KiB, 64KiB, 1MiB, 16MiB
- Size (x86)
 - PD: 4KiB, 4 byte slots
 - PT: 4KiB, 4 byte slots
 - Frame: 4KiB, 4MiB

seL4 VSpace: API



- PD
 - `seL4_<ARCH>_PageTable_Map`
 - `seL4_<ARCH>_PageTable_Unmap`
 - Can map in large frames directly:
 - `seL4_<ARCH>_Page_Map`
 - `seL4_<ARCH>_Page_Unmap`
- PT
 - `seL4_<ARCH>_Page_Map`
 - `seL4_<ARCH>_Page_Unmap`
- Note:
 - `<ARCH>` is either ARM or IA32
 - Libsel4utils provides architecture independent versions of calls
 - E.g.: `seL4_ARCH_Page_Map`

Hands-on: hello-3



- Extends code-base from hello-2 to do IPC between threads
- Prepare
 - `make ia32_hello-3_defconfig`
 - `make`
 - Edit `apps/hello-3/src/main.c`
- Root task main thread:
 - TODO 1: get a frame cap for the ipc buffer
 - TODO 2: try to map the frame the first time
 - TODO 3: create a page table
 - TODO 4: map the page table
 - TODO 5: then map the frame in
 - TODO 6: create an endpoint
 - TODO 7: make a badged copy of it
 - TODO 8: set the data to send
 - TODO 9: send & wait for reply
 - TODO 10: get reply message
- Second thread:
 - TODO 11: wait for a message
 - TODO 12-13: get & check message
 - TODO 14-15: send message back

More seL4 API and libraries

- Starting a new process**

What's a seL4 process?



- seL4 doesn't have a concept of “process”
- Traditional process:
 - Each process has its own:
 - CSpace & VSpace & Threads
- Non-traditional process:
 - Different combinations of VSpace and CSpace sharing:
 - Shared CSpace, separate VSpaces
 - Separate CSpace, shared Vspace
 - Partially shared CSpace
 - Partially shared VSpace

Build system: non-root task apps



- Adding a non-root-task app, Steps:
 - Example: root task: hello-4, non root-task app: hello-4-app
 - Make a new app for it (e.g. hello-app-4), with config and src files
 - Add dependency information about the new app in root task app's Kbuild
 - `hello-4-components-y += hello-4-app`
 - `hello-4-components = $(addprefix $(STAGE_BASE)/bin/, $(hello-4-components-y))`
 - `hello-4: export COMPONENTS=${hello-4-components}`
 - `hello-4: ${hello-4-components-y} kernel_elf $(hello-4-y)`
 - And in the root task app's Makefile to add it to the cpio archive
 - `{COMPONENTS}: false`
 - `archive.o: ${COMPONENTS} $(Q)mkdir -p $(dir $@) ${COMMON_PATH}/files_to_obj.sh $@ _cpio_archive $^`

Library: vspace



- Interface for managing VSpaces
 - manage current VSpace
 - manage other VSpaces
 - note: create is not part of vspace API !
 - allocate frames and map them into a VSpace
- Key concepts
 - **reservation**: portion of VSpace, that will not be given to others
 - **mapping**: frame mapped into a VSpace at a virtual address
- Interfaces defined
 - `vspace`
- Implemented by
 - `sel4utils`

Vspace: API



- Files
 - libseL4vspace, `#include <vspace/vspace.h>`
- Datastructs
 - `vspace_t`
 - `reservation_t`: a reserved range of `vspace` addresses
- Functions
 - `vspace_reserve_range`, `vspace_free_reservation`
 - `vspace_new_pages`: create frames and map into VSpace
 - `vspace_map_pages`: map given frames into VSpace
 - `vspace_unmap_pages`: provides different ways to free frame object
 - `vspace_get_cap`: get frame cap for specific virtual address
 - `vspace_get_root`: get cap to PD of VSpace

Library: sel4utils



- Utility code to make life easier
 - create and manage threads and processes
 - create vspaces, implement vspace interface
 - load ELF code
- Key concepts
 - process: CSpace + VSpace + TCB
- Interfaces implemented
 - `vspace`
 - `sel4utils`: util functions provided by the library
 - Thread and process management
 - Logging and profiling
 - Architecture agnostic functions

seL4utils: API



- Files:
 - libseL4utils, #include <sel4utils/...>: vspace.h, process.h, mapping.h
 - libs/libsel4utils/include/sel4utils
- Datastructs
 - process_t
 - thread_t
- Functions
 - sel4utils_bootstrap_vspace_with_bootinfo
 - sel4utils_get_vspace: **create new vspace**
 - sel4utils_configure_process: **create process**
 - sel4utils_spawn_process_v: **start process**
 - sel4utils_*_cap_to_process: **mint, copy**
 - seL4_ARCH_*: **architecture independent wrapper for seL4 syscalls**

Hands-on: hello-4 (part 1)



- Create a new process
- Prepare
 - `make ia32_hello-4_defconfig`
 - `make`
- **Edit** `apps/hello-4/src/main.c`
 - TODO 1: create a vspace object
 - TODO 2-3: use `sel4utils` to make a new process
 - TODO 6: spawn the process
- **Edit** `apps/hello-4-app/src/main.c`
 - Fix it so that the asserts don't fail
- **Build and Run**
 - New process should print something

Hands-on: hello-4 (part 2)



- Add IPC between the root task and the new process
- Edit `apps/hello-4/src/main.c`
 - TODO 4: make a cspacepath for the new endpoint cap
 - TODO 5: copy the endpoint cap and add a badge to the new cap
 - TODO 7: wait for a message
 - TODO 8: send the modified message back
- Edit `apps/hello-4-app/src/main.c`
 - TODO 9: send and wait for a reply
- Build and run
 - See message sent back and forth!

What's Next?

Advanced seL4



- In-depth VSpace and CSpace
 - See seL4 manual: <http://sel4.systems/Info/Docs/seL4-manual.pdf>
- Reimplement hello-2 (and 3 and 4) using only seL4 API
- CapDL: capability distribution language – static system setup
 - <https://github.com/seL4/capDL-tool>
 - <https://github.com/seL4/capdl-loader-app>
- Device drivers (in seL4; or with CAmkES)
 - Libplatsupport: <https://github.com/seL4/libplatsupport>
 - Libsel4platsupport: <https://github.com/seL4/libsel4platsupport>
- A whole operating system (e.g. UNSW Advanced Operating Systems Assignments)
 - <http://www.cse.unsw.edu.au/~cs9242>
 - <https://bitbucket.org/kevinelp/unsw-advanced-operating-systems>